

# Diabetes\_Predction

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import LabelEncoder
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from xgboost import XGBClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
import warnings
warnings.filterwarnings("ignore")
```

```
In [2]: df=pd.read_csv(r"C:\Users\JANHAVI\Desktop\diabetes_prediction_dataset.csv")
```

```
In [3]: df.head()
```

```
Out[3]:
```

|   | gender | age  | hypertension | heart_disease | smoking_history | bmi   | HbA1c_level | blood_glucose_ |
|---|--------|------|--------------|---------------|-----------------|-------|-------------|----------------|
| 0 | Female | 80.0 | 0            | 1             | never           | 25.19 | 6.6         |                |
| 1 | Female | 54.0 | 0            | 0             | No Info         | 27.32 | 6.6         |                |
| 2 | Male   | 28.0 | 0            | 0             | never           | 27.32 | 5.7         |                |
| 3 | Female | 36.0 | 0            | 0             | current         | 23.45 | 5.0         |                |
| 4 | Male   | 76.0 | 1            | 1             | current         | 20.14 | 4.8         |                |

```
In [4]: df.isna().any()
```

```
Out[4]:
```

|                     |       |
|---------------------|-------|
| gender              | False |
| age                 | False |
| hypertension        | False |
| heart_disease       | False |
| smoking_history     | False |
| bmi                 | False |
| HbA1c_level         | False |
| blood_glucose_level | False |
| diabetes            | False |
| dtype:              | bool  |

```
In [5]: df.corr(numeric_only=True)
```

Out[5]:

|                     | age      | hypertension | heart_disease | bmi      | HbA1c_level | blood_glucose_level |
|---------------------|----------|--------------|---------------|----------|-------------|---------------------|
| age                 | 1.000000 | 0.251171     | 0.233354      | 0.337396 | 0.101354    | 0.110               |
| hypertension        | 0.251171 | 1.000000     | 0.121262      | 0.147666 | 0.080939    | 0.084               |
| heart_disease       | 0.233354 | 0.121262     | 1.000000      | 0.061198 | 0.067589    | 0.070               |
| bmi                 | 0.337396 | 0.147666     | 0.061198      | 1.000000 | 0.082997    | 0.091               |
| HbA1c_level         | 0.101354 | 0.080939     | 0.067589      | 0.082997 | 1.000000    | 0.166               |
| blood_glucose_level | 0.110672 | 0.084429     | 0.070066      | 0.091261 | 0.166733    | 1.000               |
| diabetes            | 0.258008 | 0.197823     | 0.171727      | 0.214357 | 0.400660    | 0.419               |

In [6]: `df.shape`

Out[6]: (100000, 9)

## Unique Elements

In [7]: `for column in df.columns:  
 unique_values = df[column].unique()  
 print('Column "{}" has unique values: {}'.format(column, unique_values))`

Column "gender" has unique values: ['Female' 'Male' 'Other']  
 Column "age" has unique values: [80. 54. 28. 36. 76. 20. 44. 79. 4  
 2. 32. 53. 78.  
 67. 15. 37. 40. 5. 69. 72. 4. 30. 45. 43. 50.  
 41. 26. 34. 73. 77. 66. 29. 60. 38. 3. 57. 74.  
 19. 46. 21. 59. 27. 13. 56. 2. 7. 11. 6. 55.  
 9. 62. 47. 12. 68. 75. 22. 58. 18. 24. 17. 25.  
 0.08 33. 16. 61. 31. 8. 49. 39. 65. 14. 70. 0.56  
 48. 51. 71. 0.88 64. 63. 52. 0.16 10. 35. 23. 0.64  
 1.16 1.64 0.72 1.88 1.32 0.8 1.24 1. 1.8 0.48 1.56 1.08  
 0.24 1.4 0.4 0.32 1.72 1.48]  
 Column "hypertension" has unique values: [0 1]  
 Column "heart\_disease" has unique values: [1 0]  
 Column "smoking\_history" has unique values: ['never' 'No Info' 'current' 'former'  
 'ever' 'not current']  
 Column "bmi" has unique values: [25.19 27.32 23.45 ... 59.42 44.39 60.52]  
 Column "HbA1c\_level" has unique values: [6.6 5.7 5. 4.8 6.5 6.1 6. 5.8 3.5 6.2  
 4. 4.5 9. 7. 8.8 8.2 7.5 6.8]  
 Column "blood\_glucose\_level" has unique values: [140 80 158 155 85 200 145 100 1  
 30 160 126 159 90 260 220 300 280 240]  
 Column "diabetes" has unique values: [0 1]

In [8]: `df["smoking_history"].value_counts()`

Out[8]:

|                           |       |
|---------------------------|-------|
| smoking_history           |       |
| No Info                   | 35816 |
| never                     | 35095 |
| former                    | 9352  |
| current                   | 9286  |
| not current               | 6447  |
| ever                      | 4004  |
| Name: count, dtype: int64 |       |

In [9]: `df["smoking_history"].value_counts()/len(df)`

```
Out[9]: smoking_history
No Info      0.35816
never        0.35095
former       0.09352
current      0.09286
not current  0.06447
ever         0.04004
Name: count, dtype: float64
```

```
In [10]: # Replaceing No Info columns with pd.NA
df['smoking_history'] = df['smoking_history'].replace('No Info', pd.NA)

# Replace missing values with the mode it is string so we are using mode
mode_value = df['smoking_history'].mode()[0]
df['smoking_history'] = df['smoking_history'].fillna(mode_value) #filling no info v

# Printing the updated value counts
print(df['smoking_history'].value_counts())
```

```
smoking_history
never      70911
former     9352
current    9286
not current 6447
ever       4004
Name: count, dtype: int64
```

```
In [11]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 100000 entries, 0 to 99999
Data columns (total 9 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   gender                100000 non-null object
 1   age                  100000 non-null float64
 2   hypertension          100000 non-null int64
 3   heart_disease         100000 non-null int64
 4   smoking_history       100000 non-null object
 5   bmi                   100000 non-null float64
 6   HbA1c_level           100000 non-null float64
 7   blood_glucose_level   100000 non-null int64
 8   diabetes              100000 non-null int64
dtypes: float64(3), int64(4), object(2)
memory usage: 6.9+ MB
```

```
In [12]: df.gender.value_counts()
```

```
Out[12]: gender
Female    58552
Male      41430
Other      18
Name: count, dtype: int64
```

```
In [13]: df.describe()
```

Out[13]:

|              | age           | hypertension  | heart_disease | bmi           | HbA1c_level   | blood_glucose_l |
|--------------|---------------|---------------|---------------|---------------|---------------|-----------------|
| <b>count</b> | 100000.000000 | 100000.000000 | 100000.000000 | 100000.000000 | 100000.000000 | 100000.000      |
| <b>mean</b>  | 41.885856     | 0.07485       | 0.039420      | 27.320767     | 5.527507      | 138.058         |
| <b>std</b>   | 22.516840     | 0.26315       | 0.194593      | 6.636783      | 1.070672      | 40.708          |
| <b>min</b>   | 0.080000      | 0.00000       | 0.000000      | 10.010000     | 3.500000      | 80.000          |
| <b>25%</b>   | 24.000000     | 0.00000       | 0.000000      | 23.630000     | 4.800000      | 100.000         |
| <b>50%</b>   | 43.000000     | 0.00000       | 0.000000      | 27.320000     | 5.800000      | 140.000         |
| <b>75%</b>   | 60.000000     | 0.00000       | 0.000000      | 29.580000     | 6.200000      | 159.000         |
| <b>max</b>   | 80.000000     | 1.00000       | 1.000000      | 95.690000     | 9.000000      | 300.000         |

```
In [14]: #removing , in bmi parameter
df["bmi"] = [float(str(i).replace(",","")) for i in df["bmi"]]
```

```
In [15]: #plotting value_counts of diabetes in graphical representation
df['diabetes'].value_counts().plot(kind='barh')


#Xlabel name
plt.xlabel('count')

#ylabel name
plt.ylabel('diabetes')

#title of the plot
plt.title('count of diabetes and Non diabetes')

#invert ylabes to no diabetes on top
plt.gca().invert_yaxis()

#printing the plot
plt.show()
```



In [16]: `df['diabetes'].value_counts()/len(df) #percentage of 1--diabetes and 2--no diabetes`

Out[16]:

```
diabetes
0    0.915
1    0.085
Name: count, dtype: float64
```

In [17]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 100000 entries, 0 to 99999
Data columns (total 9 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   gender                100000 non-null object
 1   age                   100000 non-null float64
 2   hypertension          100000 non-null int64
 3   heart_disease         100000 non-null int64
 4   smoking_history       100000 non-null object
 5   bmi                   100000 non-null float64
 6   HbA1c_level           100000 non-null float64
 7   blood_glucose_level   100000 non-null int64
 8   diabetes              100000 non-null int64
dtypes: float64(3), int64(4), object(2)
memory usage: 6.9+ MB
```

In [18]: `le=LabelEncoder() #activating label encoder function`

le

Out[18]:

▼ LabelEncoder ⓘ ?

► Parameters

```
In [19]: Label_encod_columns=['gender','smoking_history'] #selecting columns to apply label
df[Label_encod_columns]=df[Label_encod_columns].apply(le.fit_transform) #applying l
```


```
In [20]: df.head(3)
```

```
Out[20]:
```

|   | gender | age  | hypertension | heart_disease | smoking_history | bmi   | HbA1c_level | blood_glucose_ |
|---|--------|------|--------------|---------------|-----------------|-------|-------------|----------------|
| 0 | 0      | 80.0 | 0            | 1             | 3               | 25.19 | 6.6         |                |
| 1 | 0      | 54.0 | 0            | 0             | 3               | 27.32 | 6.6         |                |
| 2 | 1      | 28.0 | 0            | 0             | 3               | 27.32 | 5.7         |                |


```
In [21]: sns.boxplot(data=df[['age','blood_glucose_level','bmi']]) #checking outliers using
```

```
Out[21]: <Axes: >
```




```
In [22]: sns.boxplot(data=df['HbA1c_level'])
```

```
Out[22]: <Axes: ylabel='HbA1c_level'>
```

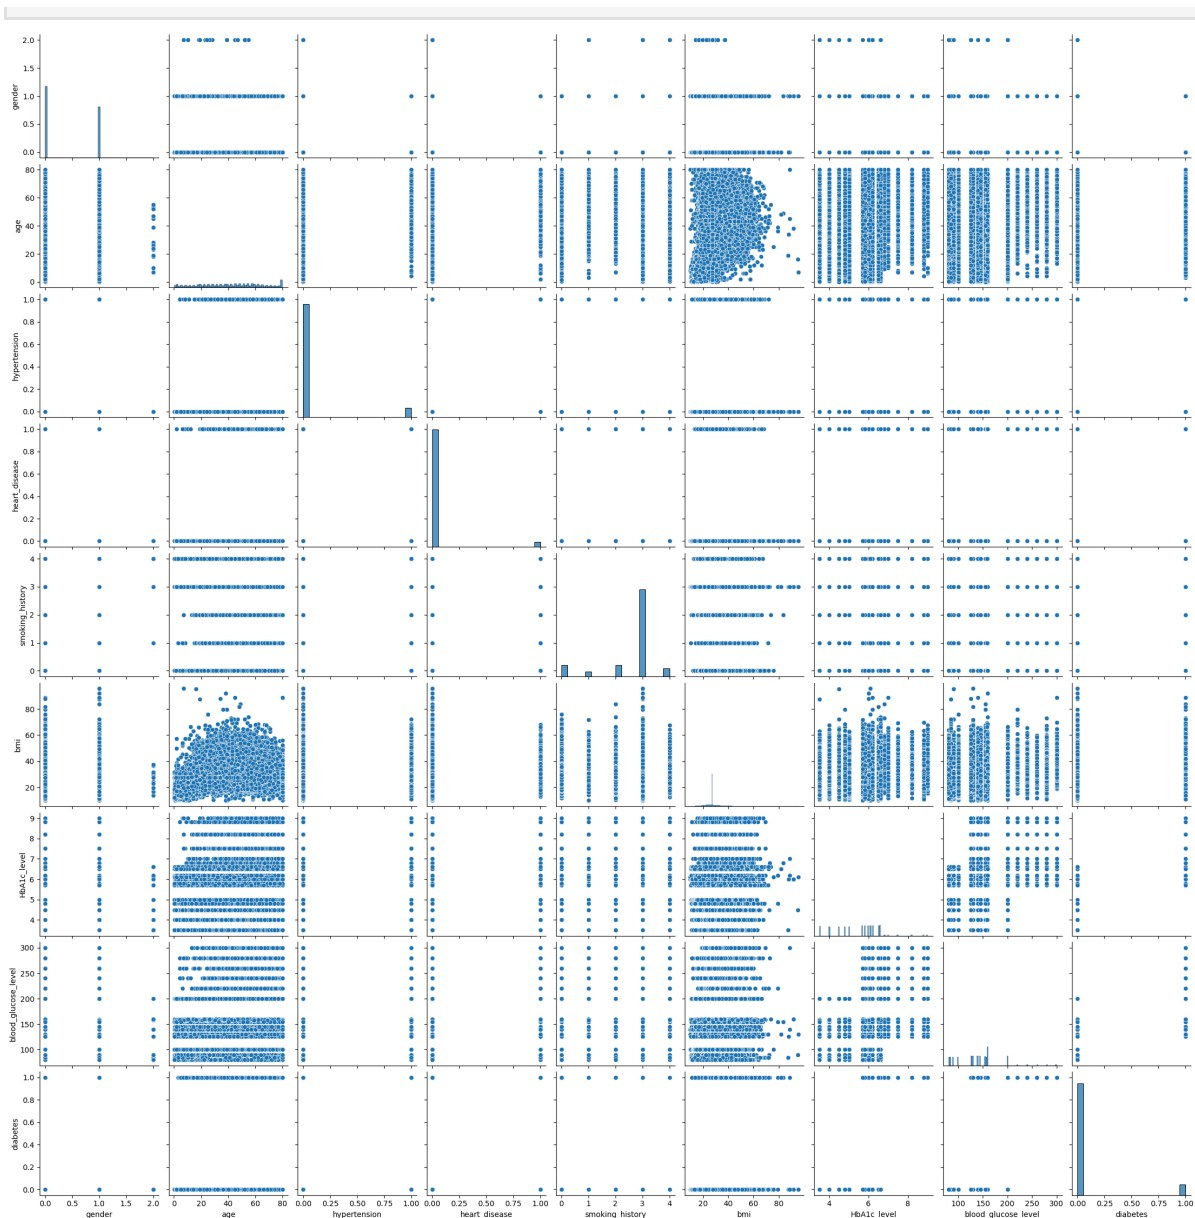


```
In [23]: sns.lmplot(data=df, x='blood_glucose_level', y='diabetes', fit_reg=False)#implot pl
```

```
Out[23]: <seaborn.axisgrid.FacetGrid at 0x23078334c90>
```



```
In [24]: sns.pairplot(df) #using pairplot to check relation between parameters  
  
#print the pairplot  
plt.show()
```



In [25]: `df.corr()`

Out[25]:


|                     | gender    | age       | hypertension | heart_disease | smoking_history | bmi       |
|---------------------|-----------|-----------|--------------|---------------|-----------------|-----------|
| gender              | 1.000000  | -0.030656 | 0.014203     | 0.077696      | -0.044081       | -0.022994 |
| age                 | -0.030656 | 1.000000  | 0.251171     | 0.233354      | -0.098969       | 0.337396  |
| hypertension        | 0.014203  | 0.251171  | 1.000000     | 0.121262      | -0.048631       | 0.147666  |
| heart_disease       | 0.077696  | 0.233354  | 0.121262     | 1.000000      | -0.048253       | 0.061198  |
| smoking_history     | -0.044081 | -0.098969 | -0.048631    | -0.048253     | 1.000000        | -0.087735 |
| bmi                 | -0.022994 | 0.337396  | 0.147666     | 0.061198      | -0.087735       | 1.000000  |
| HbA1c_level         | 0.019957  | 0.101354  | 0.080939     | 0.067589      | -0.017534       | 0.082997  |
| blood_glucose_level | 0.017199  | 0.110672  | 0.084429     | 0.070066      | -0.022985       | 0.091261  |
| diabetes            | 0.037411  | 0.258008  | 0.197823     | 0.171727      | -0.049841       | 0.214357  |

In [26]: `plt.figure(figsize=(20,8)) #figsize`



```
#printing graphical representations of
df.corr()['diabetes'].sort_values(ascending=False).plot(kind='bar')
```

Out[26]: <Axes: >



In [27]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 100000 entries, 0 to 99999
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   gender                 100000 non-null int32
1   age                   100000 non-null float64
2   hypertension           100000 non-null int64
3   heart_disease          100000 non-null int64
4   smoking_history        100000 non-null int32
5   bmi                    100000 non-null float64
6   HbA1c_level            100000 non-null float64
7   blood_glucose_level    100000 non-null int64
8   diabetes               100000 non-null int64
dtypes: float64(3), int32(2), int64(4)
memory usage: 6.1 MB
```

In [28]: #selecting X variables  
X = df.loc[:, 'age':'heart\_disease'].join(df.loc[:, 'bmi':'blood\_glucose\_level'])

X

Out[28]:

|       | age  | hypertension | heart_disease | bmi   | HbA1c_level | blood_glucose_level |
|-------|------|--------------|---------------|-------|-------------|---------------------|
| 0     | 80.0 | 0            | 1             | 25.19 | 6.6         | 140                 |
| 1     | 54.0 | 0            | 0             | 27.32 | 6.6         | 80                  |
| 2     | 28.0 | 0            | 0             | 27.32 | 5.7         | 158                 |
| 3     | 36.0 | 0            | 0             | 23.45 | 5.0         | 155                 |
| 4     | 76.0 | 1            | 1             | 20.14 | 4.8         | 155                 |
| ...   | ...  | ...          | ...           | ...   | ...         | ...                 |
| 99995 | 80.0 | 0            | 0             | 27.32 | 6.2         | 90                  |
| 99996 | 2.0  | 0            | 0             | 17.37 | 6.5         | 100                 |
| 99997 | 66.0 | 0            | 0             | 27.83 | 5.7         | 155                 |
| 99998 | 24.0 | 0            | 0             | 35.42 | 4.0         | 100                 |
| 99999 | 57.0 | 0            | 0             | 22.43 | 6.6         | 90                  |

100000 rows × 6 columns

```
In [29]: y=df.loc[:, 'diabetes'] #y variable
y #printing y variable
```

```
Out[29]: 0      0
1      0
2      0
3      0
4      0
..
99995  0
99996  0
99997  0
99998  0
99999  0
Name: diabetes, Length: 100000, dtype: int64
```

```
In [30]: # splitting trining and testing data in 70 30 rating testing size is 0.3 random_stat
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.2,random_state=0)
```

```
In [31]: X_train.head()
```

```
Out[31]:
```

|       | age  | hypertension | heart_disease | bmi   | HbA1c_level | blood_glucose_level |
|-------|------|--------------|---------------|-------|-------------|---------------------|
| 10382 | 2.0  | 0            | 0             | 16.45 | 6.2         | 159                 |
| 73171 | 55.0 | 0            | 0             | 24.59 | 6.0         | 130                 |
| 30938 | 24.0 | 0            | 0             | 21.77 | 4.5         | 130                 |
| 99310 | 30.0 | 0            | 0             | 27.32 | 6.2         | 159                 |
| 58959 | 13.0 | 0            | 0             | 18.37 | 6.5         | 130                 |

```
In [32]: print('Shape of Train data')
print(X_train.shape)
```

```
print(y_train.shape)

print('Shape of Testing data')

print(X_test.shape)

print(y_test.shape)
```

```
Shape of Train data
(80000, 6)
(80000,)
Shape of Testing data
(20000, 6)
(20000,)
```

```
In [33]: ss=StandardScaler() #activating StandardScaler()

ss
```

```
Out[33]: ▾ StandardScaler ⓘ ?
          ► Parameters
```

```
In [34]: X_train_scaled=ss.fit_transform(X_train) #scaling X_train data
```

```
In [35]: if len(X_test.shape) == 1:    #if x is 1d array
          X_test = X_test.values.reshape(-1, 1) #converting to 2d array

          X_test_scaled = ss.fit_transform(X_test) #scaling X_test data
```

```
In [36]: model_lr=LogisticRegression() #activating Logistic Regression
```

```
In [37]: model_lr.fit(X_train_scaled,y_train) #training logistic regression model
```

```
Out[37]: ▾ LogisticRegression ⓘ ?
          ► Parameters
```

```
In [38]: y_pred=model_lr.predict(X_test_scaled) #predecting y_test data
          y_pred[:10]
```

```
Out[38]: array([0, 0, 0, 0, 0, 0, 0, 0, 0, 0], dtype=int64)
```

```
In [39]: y_test[:10] # actual y_test data
```

```
Out[39]: 3582      0
          60498    0
          53227    0
          21333    0
          3885     0
          51521    0
          84261    0
          10685    1
          59948    0
          41032    0
          Name: diabetes, dtype: int64
```

```
In [40]: accuracy_score(y_pred,y_test) #accuracy_score
```

Out[40]: 0.95975

In [41]: `print(classification_report(y_pred,y_test)) #classifiaction_report`

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.99      | 0.97   | 0.98     | 18736   |
| 1            | 0.63      | 0.86   | 0.73     | 1264    |
| accuracy     |           |        | 0.96     | 20000   |
| macro avg    | 0.81      | 0.91   | 0.85     | 20000   |
| weighted avg | 0.97      | 0.96   | 0.96     | 20000   |


In [42]: `confusion_matrix(y_pred,y_test) #confusion_matrix`

Out[42]: `array([[18114, 622],  
[ 183, 1081]], dtype=int64)`

In [43]: `y_train.value_counts() #data is highly imblancing`

Out[43]: `diabetes  
0 73203  
1 6797  
Name: count, dtype: int64`

In [44]: `value_counts=y_train.value_counts()  
  
plt.figure(figsize=(16, 8))  
  
plt.pie(value_counts, labels=value_counts.index, autopct='%1.2f%%', startangle=140)  
  
plt.title('Distribution of y_train')  
  
plt.axis('equal') # Equal aspect ratio ensures that pie is drawn as a circle.  
  
plt.show()`



In [45]: `from imblearn.over_sampling import SMOTE # using smote function to balance our set  
  
smote=SMOTE()  
  
X_ovs,y_ovs=smote.fit_resample(X,y) #passing X and y variables to it to balance out`


```
fig, oversp = plt.subplots()

oversp.pie( y_ovs.value_counts(), autopct='%.2f')

oversp.set_title("Over-sampling")

plt.show()
```

Over-sampling



```
In [46]: # Dividing our resampling data into 70 30 ratio

Xr_train,Xr_test,yr_train,yr_test=train_test_split(X_ovs,y_ovs,train_size=0.7,randc
```

```
In [47]: print('train data shape')

print(Xr_train.shape)

print(yr_train.shape)

print('test data shape')

print(Xr_test.shape)

print(yr_test.shape)
```

```
train data shape
(128099, 6)
(128099,)
test data shape
(54901, 6)
(54901,)
```

```
In [48]: print('y_train and y_test value_count')
print(yr_train.value_counts())
print(yr_test.value_counts())
```

```

y_train and y_test value_count
diabetes
0    64131
1    63968
Name: count, dtype: int64
diabetes
1    27532
0    27369
Name: count, dtype: int64

```

In [49]: `ss=StandardScaler()`

ss

Out[49]:

▼ StandardScaler ⓘ ?

► Parameters

In [50]: `data=Xr_train,Xr_test`

```

xr_train_sc=ss.fit_transform(Xr_train) # scaling our resampling data xr train

Xr_test_sc=ss.fit_transform(Xr_test) # scaling our resamplig xr_test data

```

In [51]: `Xr_train_scaled = pd.DataFrame(xr_train_sc) #Xr_train_scaled converting into the dataframe`

```

print(Xr_train_scaled.shape)
Xr_train_scaled.head()
print(yr_train.shape)

```

```

(128099, 6)
(128099,)

```

In [52]: `Xr_test_scaled=pd.DataFrame(Xr_test_sc) #Xr_test converting into the dataframe`

```

print(Xr_test_scaled.shape)
Xr_test_scaled.head()

```

```

(54901, 6)

```

Out[52]:

|   | 0         | 1         | 2         | 3         | 4         | 5         |
|---|-----------|-----------|-----------|-----------|-----------|-----------|
| 0 | 0.957108  | -0.294735 | -0.204364 | 0.539410  | -0.111145 | 2.029416  |
| 1 | -1.095860 | -0.294735 | -0.204364 | -0.404865 | 0.370604  | -0.064547 |
| 2 | -1.467180 | -0.294735 | -0.204364 | -0.287066 | 0.370604  | -1.460522 |
| 3 | -0.770954 | 3.392878  | -0.204364 | 0.293712  | 0.370604  | -1.373274 |
| 4 | -1.374350 | -0.294735 | -0.204364 | -0.287066 | -2.150914 | -1.111528 |

In [53]: `model_lk=LogisticRegression()`

```

model_lk.fit(Xr_train_scaled,yr_train) #training the model

```

Out[53]:

▼
LogisticRegression

Parameters

```
In [54]: y_pred_lr=model_lr.predict(Xr_test_scaled) #predicting yr_test data
y_pred_lr[:10]
```

```
Out[54]: array([1, 0, 0, 0, 0, 1, 0, 0, 0, 1], dtype=int64)
```

```
In [55]: yr_test[:10]
```

```
Out[55]: 180328    1
573         0
13494       0
93981       0
75389       0
180973      1
71021       0
19293       0
16393       0
121419      1
Name: diabetes, dtype: int64
```

```
In [56]: #classification_report for predict value and original value
print(classification_report(y_pred_lr,yr_test))
```

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.88      | 0.88   | 0.88     | 27383   |
| 1            | 0.88      | 0.88   | 0.88     | 27518   |
| accuracy     |           |        | 0.88     | 54901   |
| macro avg    | 0.88      | 0.88   | 0.88     | 54901   |
| weighted avg | 0.88      | 0.88   | 0.88     | 54901   |

```
In [57]: #confusion_matrix for predict value and original value
confusion_matrix(y_pred_lr,yr_test)
```

```
Out[57]: array([[24174,  3209],
               [ 3195, 24323]], dtype=int64)
```

## Decision Tree Classifier

```
In [58]: # activating DecisionTree Classifier
model_dtc=DecisionTreeClassifier()

# passing xr_train_scaled, yr_train to training the model
model_dtc.fit(Xr_train_scaled,yr_train)

model_dtc
```

Out[58]:

▼
DecisionTreeClassifier

Parameters

```
In [59]: y_pred_dtc=model_dtc.predict(Xr_test_scaled) # predicting yr_test data
```

```
In [60]: # classification report for decisionTreeClassifier

print(classification_report(y_pred_dtc,yr_test))
```

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.64      | 1.00   | 0.78     | 17500   |
| 1            | 1.00      | 0.73   | 0.85     | 37401   |
| accuracy     |           |        | 0.82     | 54901   |
| macro avg    | 0.82      | 0.87   | 0.81     | 54901   |
| weighted avg | 0.88      | 0.82   | 0.82     | 54901   |

```
In [61]: confusion_matrix(y_pred_dtc,yr_test)
```

```
Out[61]: array([[17433, 67],
               [ 9936, 27465]], dtype=int64)
```

## Random Forest Classifier

```
In [62]: model_rfc=RandomForestClassifier() #activating the fuction

model_rfc.fit(Xr_train_scaled,yr_train)
```

```
Out[62]: ▼ RandomForestClassifier ⓘ ?
          ► Parameters
```

```
In [63]: y_pred_rfc=model_rfc.predict(Xr_test_scaled)
```

```
In [64]: print(classification_report(y_pred_rfc,yr_test))
```

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.78      | 0.99   | 0.87     | 21526   |
| 1            | 0.99      | 0.82   | 0.90     | 33375   |
| accuracy     |           |        | 0.89     | 54901   |
| macro avg    | 0.89      | 0.90   | 0.88     | 54901   |
| weighted avg | 0.91      | 0.89   | 0.89     | 54901   |

```
In [65]: confusion_matrix(y_pred_rfc,yr_test)
```

```
Out[65]: array([[21306, 220],
               [ 6063, 27312]], dtype=int64)
```


## XGBOOST

```
In [66]: model_xgb=XGBClassifier()

model_xgb.fit(Xr_train_scaled,yr_train)
```



Out[66]:

In [67]: `y_pred_xgb=model_xgb.predict(Xr_test_scaled)`In [68]: `print(classification_report(y_pred_xgb,yr_test))`

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.89      | 0.96   | 0.92     | 25342   |
| 1            | 0.96      | 0.90   | 0.93     | 29559   |
| accuracy     |           |        | 0.93     | 54901   |
| macro avg    | 0.93      | 0.93   | 0.93     | 54901   |
| weighted avg | 0.93      | 0.93   | 0.93     | 54901   |

In [69]: `confusion_matrix(y_pred_xgb,yr_test)`Out[69]: `array([[24305, 1037],  
 [ 3064, 26495]], dtype=int64)`

## Hyperparameter Tunning & Best Parameter GridSearchcv

```
In [ ]: from sklearn.model_selection import GridSearchCV, cross_val_score
from sklearn.linear_model import LogisticRegression

# Define the parameter grid to search over
param_grid = {
    'C': [0.001, 0.01, 0.1, 1, 10, 100], # Regularization parameter
    'penalty': ['l1', 'l2'] # Penalty type
}

# Create a Logistic Regression model
logistic = LogisticRegression()

# Create a GridSearchCV object
grid_search = GridSearchCV(estimator=logistic, param_grid=param_grid, cv=10)

# Initialize an empty list to store the accuracy scores
accuracy_scores = []

# Perform cross-validation 10 times
for _ in range(10):
    # Fit the GridSearchCV object to the training data
    grid_search.fit(Xr_train_scaled, yr_train)

    # Get the best parameters
    best_params = grid_search.best_params_

    # Perform cross-validation with the best model
    cv_scores = cross_val_score(grid_search.best_estimator_, Xr_train_scaled, yr_train)

    # Store the mean accuracy score
    accuracy_scores.append(cv_scores.mean())

# Print the accuracy scores obtained over 10 iterations
```

```
#print("Accuracy scores over 10 iterations:", accuracy_scores)
print("Accuracy scores over 10 iterations:", [{":.2f}".format(score) for score in accuracy_scores])

# Get the best parameters and best score
best_params = grid_search.best_params_
best_score = grid_search.best_score_

print("Best parameters found:", best_params)
print("Best cross-validation score:", best_score)
```

## Final Model

```
In [74]: from sklearn.linear_model import LogisticRegression

# Create a Logistic Regression model with the best parameters
final_model = LogisticRegression(C=0.001, penalty='l2')

# Fit the final model to the entire training dataset
final_model.fit(Xr_train_scaled, yr_train)
```

```
Out[74]: LogisticRegression
Parameters
```

```
In [75]: import pickle

# Save the final model to a pickle file
with open('final_model.pkl', 'wb') as file:
    pickle.dump(final_model, file)
```

```
In [76]: import pickle
import numpy as np

# Load the model from the pickle file
with open('final_model.pkl', 'rb') as file:
    loaded_model = pickle.load(file)

# Define the mean and standard deviation of the training data
mean_values = [41.885856, 0.07485, 0.03942, 27.320767, 5.527507, 138.058060]
std_values = [22.516840, 0.26315, 0.194593, 6.636783, 1.070672, 40.708136]

# Define the input features for prediction
age = 30
hypertension = 0
heart_disease = 0
bmi = 100.0
HbA1c_level = 5.0
blood_glucose_level = 90

# Scale the input features manually
scaled_features = [(x - mean) / std for x, mean, std in zip(
    [age, hypertension, heart_disease, bmi, HbA1c_level, blood_glucose_level],
    mean_values, std_values
)]

# Make predictions on the scaled data
prediction = loaded_model.predict([scaled_features])
```

```
# Print the prediction
if prediction[0] == 1:
    print("Diabetic")
else:
    print("Not Diabetic")
```

Diabetic

In [ ]: